

上海科技大学
2024-2025 强化学习应用实践
Project1 Part-A

姓名: 周守琛

学号: 2021533042

2025 年 3 月 17 日

- 状态集合 S : $S = \{s_1, s_2, s_3\}$

- 动作集合 A : $A = \{a_1, a_2\}$

- 终止状态: s_3 为终止状态, 不再执行动作。

- 转移概率 $P(s'|s, a)$:

– 从 s_1 :

$$P(s_2|s_1, a_1) = 0.5, \quad P(s_3|s_1, a_1) = 0.5$$

$$P(s_2|s_1, a_2) = 0.7, \quad P(s_3|s_1, a_2) = 0.3$$

– 从 s_2 :

$$P(s_1|s_2, a_1) = 0.6, \quad P(s_3|s_2, a_1) = 0.4$$

$$P(s_1|s_2, a_2) = 0.8, \quad P(s_3|s_2, a_2) = 0.2$$

– 从 s_3 : 无后续转移 (所有 $P(s'|s_3, a) = 0$)。

- 奖励函数 $R(s, a)$:

$$R(s_1, a_1) = 1, \quad R(s_1, a_2) = 2$$

$$R(s_2, a_1) = 3, \quad R(s_2, a_2) = 0$$

$$R(s_3, a) = 0 \quad \forall a$$

- 折扣因子 γ : 0.9

- 初始策略 $\pi(s, a) = \frac{1}{|A|}$: 所有非终止状态的动作选择均匀随机:

$$\pi(s_1, a_1) = \pi(s_1, a_2) = 0.5, \quad \text{其余同理。}$$

- 初始值函数/状态-动作值函数:

$$V(s) = 0 \quad \forall s, \quad Q(s, a) = 0 \quad \forall s, a$$

- 学习率 $\alpha = 0.1$ 。

1. 策略迭代 (20 分) 根据初始策略 $\pi(s, a) = \frac{1}{2}$, 手动推导策略迭代第一轮迭代步骤, 包括策略评估和策略改进, 写出:

1. 第一轮策略评估的状态值函数 $V(s)$ 。(10 分)

2. 改进后的新策略 $\pi'(s)$ 。(10 分)

Solution

We can get the value function's update rule:

$$V(s) = \sum_a \pi(s, a) \left[R(s, a) + \gamma \sum_{s'} P(s'|s, a) V(s') \right]$$

1. To get the value function $V(s)$ after the first policy evaluation, apply the formula to the initial policy, we have

$$\begin{aligned} V(s_1) &\leftarrow \frac{1}{2} \underbrace{[1 + 0.9 \times (0.5 \times 0 + 0.5 \times 0)]}_{a=a_1} + \frac{1}{2} \underbrace{[2 + 0.9 \times (0.7 \times 0 + 0.3 \times 0)]}_{a=a_2} = 1.5 \\ V(s_2) &\leftarrow \frac{1}{2} \underbrace{[3 + 0.9 \times (0.6 \times 0 + 0.4 \times 0)]}_{a=a_1} + \frac{1}{2} \underbrace{[0 + 0.9 \times (0.8 \times 0 + 0.2 \times 0)]}_{a=a_2} = 1.5 \\ V(s_3) &\leftarrow 0 \end{aligned}$$

2. To update the policy, we have the update rule

$$\pi(s) = \arg \max_a Q(s, a) = \arg \max_a \sum_{s'} P(s'|s, a) [R(s, a) + \gamma V(s')] = \arg \max_a R(s, a) + \gamma \sum_{s'} P(s'|s, a) V(s')$$

- $s = s_1, a = a_1$: $\sum_{s'} P(s'|s_1, a_1) V(s') = 1 + 0.9 \times \left[\underbrace{0.5 \times 1.5}_{s'=s_2} + \underbrace{0.5 \times 0}_{s'=s_3} \right] = 1.675$
- $s = s_1, a = a_2$: $\sum_{s'} P(s'|s_1, a_2) V(s') = 2 + 0.9 \times \left[\underbrace{0.7 \times 1.5}_{s'=s_1} + \underbrace{0.3 \times 0}_{s'=s_3} \right] = 2.945$
- $s = s_2, a = a_1$: $\sum_{s'} P(s'|s_2, a_1) V(s') = 3 + 0.9 \times \left[\underbrace{0.6 \times 1.5}_{s'=s_1} + \underbrace{0.4 \times 0}_{s'=s_3} \right] = 3.81$
- $s = s_2, a = a_2$: $\sum_{s'} P(s'|s_2, a_2) V(s') = 0 + 0.9 \times \left[\underbrace{0.8 \times 1.5}_{s'=s_1} + \underbrace{0.2 \times 0}_{s'=s_3} \right] = 1.08$

So the updated policy after the first policy improvement is

$$\pi'(s_1) = a_2, \quad \pi'(s_2) = a_1$$

Since s_3 is the terminal state, so it has no policy.

2. 价值迭代 (20 分) 根据给定的环境和初始值函数 $V(s) = 0 \forall s$, 推导价值迭代第一轮迭代步骤, 写出:

1. 每个状态的值函数 $V(s)$ 。(10 分)

2. 相应的策略 $\pi(s)$ 。(10 分)

Solution

We can get the value function's update rule:

$$V(s) = \max_a \left(R(s, a) + \gamma \sum_{s'} P(s'|s, a) V(s') \right)$$

1. To get the value function $V(s)$ after the first policy evaluation, apply the formula to the initial policy, we have

$$\begin{aligned} V(s_1) &\leftarrow \max \left(\underbrace{1 + 0.9 \times (0.5 \times 0 + 0.5 \times 0)}_{a=a_1}, \underbrace{2 + 0.9 \times (0.7 \times 0 + 0.3 \times 0)}_{a=a_2} \right) = \max(1, 2) = 2 \\ V(s_2) &\leftarrow \max \left(\underbrace{3 + 0.9 \times (0.6 \times 0 + 0.4 \times 0)}_{a=a_1}, \underbrace{0 + 0.9 \times (0.8 \times 0 + 0.2 \times 0)}_{a=a_2} \right) = \max(3, 0) = 3 \\ V(s_3) &\leftarrow 0 \end{aligned}$$

2. The update policy is

$$\pi(s) = \arg \max_a V(s)$$

So the updated policy after the first policy improvement is

$$\pi(s_1) = a_2, \quad \pi(s_2) = a_1$$